

程序代写代做 CS编程辅导

Intro to Image Understanding (CSC420)

Projects

Posted: March 12, 2024

Submission Deadline : April 19, 11.59pm, 2024

Max points for



Projects can be done individually or in pairs. If a project is done in a pair, each student will still need to defend their part. From the report it should be clear what each student contributed to the project. By the deadline you will need to hand in the project report including code. Make the code organized and documented, including scripts that run your pipeline. In the oral presentation we may ask you to run your code and be able to answer questions about your code, as well as answer questions about class material.

The oral presentation will be scheduled shortly after the project report submission deadline. The grade will evaluate the project report and an oral presentation, as well as an oral exam about class material, at the total of 40% of the final grade.

Whenever you use ideas, code or data from a paper, forum, webpage, etc, you need to cite it in your report. Whenever you use code available from some paper or method, you need to include a short description of the method showing your understanding of the technique. If there were multiple options for techniques or code, please also explain why you chose a particular one.

The grade will take into account the following factors:

- Your ideas to tackle a problem: how appropriate the techniques you chose are for the problem. Coming up with novel ideas is obviously a big plus.
- Your implementation: the accuracy of the solution, speed, and partly also how organized the code is (scripts that run the full pipeline or specific subtasks, documentation). The grade will take into account also the performance of your solution with respect to other students' results.
- Whether you implemented a technique yourself or used code available online
- How you present your results in the report: ideally you would show your results for each task by including a few pictures in your report. Even more ideally, you would show good cases where your method works and also bad cases where it doesn't. Providing some intuitive explanation of the failure cases is a plus.
- Thoroughness of your report: How well you describe the problem and the techniques you chose, how well you analyzed your results and whether your citations to the techniques you used are appropriate.
- Correctness of your answers about class material.

Project 1

A maximum of two students can work together for this project.

This project aims to explore the power of visual storytelling by developing a deep learning model that predicts the genre and budget of a movie based on a small set of random frames. This project seeks to identify visual cues that may reflect a movie's overall style, narrative, and production quality.

IMDB web scraper notebook: <https://colab.research.google.com/drive/1U6o0VpLkE-q7vBnFrSLY9K8Huuou1yDZ?usp=sharing>



- **Data Collection:** You can utilize the provided IMDB web scraping notebook to gather your training data, including movie genres, budgets, and links to their trailers (training images). Select movies from at least 10 movie genres.
- **Shot detection:** Detect shots in each movie trailer. A shot is a set of consecutive frames with a smooth camera motion.
- **Frame Extraction:** Write code to extract random frames from each movie trailer, one randomly selected frame per shot.
- **Model Development:** Design and implement a neural network (of at least three layers) that uses the extracted frames to predict the movie's genre and budget category. If two students are working on this project, please experiment with different architectures and pre-trained models to find the most effective solution. You will find the tutorial notebooks helpful here.
- **Budget Range Prediction:** Instead of a specific budget amount, predict the budget range or bucket (e.g., low, medium, and high, or, under 1 million USD, 1-10 million USD, 10-25 million USD, 25-50 million USD, and so on) to account for the variability and complexity of movie production costs.
- **Evaluation and Refinement:** Evaluate the model's performance on a separate test set.

Deliverables:

- A detailed report documenting the model architecture, training process, evaluation, and insights gained from the project. Outline key findings, challenges encountered, and potential future work.
- An Interactive Python notebook or web application demonstrating your approach (shot detection, the model's ability to predict a movie's genre and budget.)

Project 2

A maximum of two students can work together for this project.

Autonomous driving is one of the major research venues these days. A lot of effort is devoted to it by both the academics as well as industry. In this project you'll familiarize yourself with some of the most interesting problems that arise in the field of autonomous driving.

The input to your project is a stereo image pair and the camera parameters. You will also have available images where the cars have been annotated with 2D bounding boxes as well as the ground plane. Furthermore, you'll have a few images where the road has been annotated. Your task is to solve:

1. Compute disparity from the stereo images. We do not mind if you use existing code as long as you include a description of the algorithm you used, showing you understand what it is doing.
2. Compute depth of each pixel. Compute 3D location of each pixel.
3. Train a road classifier on a set of annotated images, and compute road pixels in your image. Which features would you use? Try to use both 2D and 3D features.
4. Fit a plane in 3D to the road pixels by using the depth of the pixels. Make sure your algorithm is robust to outliers.
5. Plot each pixel in 3D (we call this a 3D point cloud). On the same plot, show also the estimated ground plane.
6. Detect cars in the image. You can use the pre-trained models available here: http://kitti.is.tue.mpg.de/kitti/models_lsvm.zip, and detection code available here: <http://www.cs.berkeley.edu/~rbg/latent/>. Alternatively, you can use other detectors available online such as SSD, Fast(er)-RCNN, etc.
7. Train a classifier that predicts viewpoint for each car. The viewpoint labels are in 30° increments, thus train a 12-class classifier. Which features would you use?
8. Show a test image with the detected car bounding boxes and show the estimated viewpoints by plotting an arrow in the appropriate direction.
9. Given the ground plane, estimated depth, and the location of the car's bounding box, how would you compute a 3D bounding box around each detected car? Add the 3D bounding boxes to your plot from 5.

Deliverables:

- A detailed report documenting the method and answer to each of the questions. Outline key findings, challenges encountered, and potential future work.
- An Interactive Python notebook or web application demonstrating your method.

Project 3

This is a project for a single student (no pair work allowed).

In this project the task is to do pixel-level labeling with a set of semantic classes. In particular, you will use a dataset from the following paper:

Kota Yamaguchi, Luis E Ortiz, Tamara L Berg, "Parsing Clothing in Fashion Photographs", http://vision.cs.stanford.edu/yamaguchi/research/clothing_parsing/

The tasks for this project are:

1. (Automatically) label each pixel as either "person" or "background".
2. (Automatically) label each pixel as either "background", "skin", "hair", "t-shirt", "shoes", "pants", "dress".

You are given a train, validation and test split, as well as an evaluation function. In your report please include examples of your results as well as report performance on both tasks.

Note that the paper above provides code. While you can build on top of the paper's ideas we expect you to implement some of your own ideas as well, using newer techniques/classifiers.

Deliverables:

- A detailed report documenting the model architecture, training process, evaluation, and insights gained from the project. Outline key findings, challenges encountered, and potential future work.
- An Interactive Python notebook or web application demonstrating your approach (shot detection, the model's ability to predict a movie's genre and budget.)

程序代写代做 CS编程辅导



WeChat: cstutores

Assignment Project Exam Help

Email: tutorcs@163.com

QQ: 749389476

https://tutorcs.com

Project 4

This is a project for a single student (no pair work allowed).

Imagine a phone app where you take a picture of a book and the app tells you which book it is. In this project you will localize and recognize a book cover in an input image given a database of book covers. You will need to implement the following paper: Nister, Stewenius, *Localizing and Recognizing Book Covers with a Vocabulary Tree*, CVPR 2006, http://www-inst.eecs.berkeley.edu/~cs168/papers/nister_stewenius_cvpr2006.pdf. Here are the tasks:



1. Take a few photos of books that you have at your disposal (if you don't own a book, you can go to a library). Put the books on the table or the floor, and take photos from different angles. We will refer to these photos as *test images*. Take ten different photos:

- In the first three photos, there shall be one book visible under varying viewpoints: from more fronto-parallel to more and more oblique viewpoints.
- In the second three photos, there shall be three books visible under oblique angles. Each photo should contain different books.
- In the last 4 photos, there shall be five books visible. One of the books is occluding another book (ie one book is only half visible). In each photo create a different occlusion pattern.

2. Crawl at least 100 book cover images from the web. Also download images of covers of books from above. For example, Amazon has a wide selection of covers.
3. Implement homography estimation with RANSAC to match one book cover image to a test image.
4. Implement the efficient retrieval approach by Nister and Stewenius.
5. For a test image, localize all the books by matching them with the books from the large database. Please describe your algorithm.
6. Plot each test image with all the localized book covers and as well as the best retrieved book cover for each of them.

Deliverables:

- A detailed report documenting the method and answer to each of the questions. Outline key findings, challenges encountered, and potential future work.
- An Interactive Python notebook or web application demonstrating your method.

Project 5

This is a project for a single student (no pair work allowed).

In this project, you will use the latest 3D deep learning technology. Please read the paper <https://arxiv.org/abs/2003.01244>. Code for this work can be found here <https://github.com/bmild/nerf>. The project is to reconstruct 3D objects from multiview photographs and perform novel view synthesis.

1. Take multiview photographs of selected objects, but moving the camera around the object. The project utilizes 40 to 60 photos per object.
2. Compute camera matrices for all the photos. Please try to write your own implementation. Please also check the nerf github repository for instructions on using existing packages. Please compare your own camera estimates with the ones using existing packages.
3. Implement nerf by yourself.
4. Provide visualizations of novel view synthesis results.

Deliverables:

- A detailed report documenting the method and answer to each of the questions. Outline key findings, challenges encountered, and potential future work.
- An Interactive Python notebook or web application demonstrating your method.

程序代写代做 CS编程辅导



WeChat: cstutorcs

Assignment Project Exam Help

Email: tutors@163.com

QQ: 749389476

<https://tutorcs.com>

Useful Code and Techniques

Shot Detection. A simple way of detecting shot boundaries is to look at differences in color histograms between two consecutive frames. An even better way is to look at Displaced Frame Distances, described in Section 2.1 of this paper:

Makarand Tapaswi and Rainer Stiefelhagen, “Knock! Knock! Who is it” *Probabilistic Person Re-identification Series*, CVPR 2012, https://cvhci.anthropomatik.kit.edu/~mtapaswi/papers/knock_knock_who_is_it.pdf

More options are available in this paper:
Y. Yusoff, W. Chrihan, and A. Fidler, “A Study on Automatic Shot Change Detection. *Multimedia Applications*, 2018, <http://www.cs.utoronto.ca/~fidler/slides/CSC420/papers/shots.pdf>.

Face Detection. A few options:

- The first paper that did great on faces was Viola-Jones face detector:
Paul Viola and Michael Jones, Rapid object detection using a boosted cascade of simple features, CVPR, 2001, <https://www.cs.cmu.edu/~efros/courses/LBMV01/Papers/viola-cvpr-01.pdf>. There is lots of code online.
- Paper:
P. Felzenszwalb, R. Girshick, D. McAllester, D. Ramanan, *Object Detection with Discriminatively Trained Part-Based Models* IEEE Transactions on Pattern Analysis and Machine Intelligence, Vol. 32, No. 9, Sep. 2010 <http://cs.brown.edu/~pff/papers/lsvm-pami.pdf>
Code: <http://www.rossgirshick.info/latent/>
with a trained model called `dpm_baseline.mat` available here:
https://bitbucket.org/rodrigob/doppia/src/tip/data/trained_models/face_detection/?at=preparing_v2. This detector may work a little better. You can also check a detailed analysis and code for different face detectors here: http://markusmathias.bitbucket.org/2014_eccv_face_detection/
- Paper:
X. Zhu, D. Ramanan. *Face detection, pose estimation and landmark localization in the wild*, CVPR 2012, <http://www.ics.uci.edu/~xzhu/face/>.
This detector however also gives you the keypoints for the facial landmarks.

Tracking Here are some options (papers), choose the one it suits you best:

- Tracking via dynamic programming: <https://engineering.purdue.edu/~qobi/papers/acs2012b.pdf>
- Tracking via Hungarian method: <http://www.cvlibs.net/publications/Geiger2014PAMI.pdf>, first 3 paragraphs of Section 4.1
- Tracking via Kalman filter: lots of tutorials online

Available code:

- Check <http://www.ics.uci.edu/~dramanan/> for paper (and code): ‘H. Pirsiavash, D. Ramanan, C. Szwed, “Globally-Optimal Greedy Algorithms for Tracking a Variable Number of Objects”, Computer Vision and Pattern Recognition (CVPR), 2011

- <http://www.cvlibs.net/software/trackbydet/>

- <http://research.microsoft.com/en-us/projects/cttracking/>

Stereo For autonomous driving, it's best to check the stereo challenge of the KITTI dataset (<http://www.cvlibs.net/datasets/kitti/>). It has links to code and gives you running time of each algorithm.

- <http://user.math.uzh.ch/spangenberg/>

- <http://ttic.uchicago.edu/~dmcalister/SPS/index.html>

Flow For autonomous driving, it's best to check the stereo challenge of the KITTI dataset (<http://www.cvlibs.net/datasets/kitti/>). It has links to code and gives you running time of each algorithm.

- <http://people.csail.mit.edu/celiu/OpticalFlow/>

- Matlab has some functions to compute flow: <http://www.mathworks.com/help/vision/ref/opticalflow.html>

- OpenCV has flow and trackers, and other useful stuff: <http://opencv.org/documentation.html>

Other For features, super pixels, classifiers, etc, look under Resources of class webpage: <http://www.cs.utoronto.ca/~fidler/teaching/2019/CSC420.html>

<https://tutorcs.com>